

---

# Applied Machine Learning

## Fine-tuning the TTT Model

---

### Authors

Mihika Jain, Anya Ross, Meili Gupta  
Cornell University - CS 5875

mj434@cornell.edu, ar2535@cornell.edu, mwg76@cornell.edu

## 1 Motivation

For this project, we will be applying a pre-existing model to novel datasets and further fine-tuning the model to enhance performance. Specifically, we will be training and evaluating the Test-Time-Training (TTT) model on the Project Gutenberg, PubMed, and GovReport datasets.

Recurrent neural networks (RNNs) have long been a staple in natural language processing, yet they struggle with the vanishing gradient problem, especially when handling long input sequences. Addressing this limitation, Sun et al. recently introduced an innovative approach, "Learning to (Learn at Test Time): RNNs with Expressive Hidden States," which replaces traditional RNN hidden states with a multi-layer perceptron (MLP) to retain meaningful gradients across extended context lengths. This architecture, incorporating a concept called Test-Time Training (TTT), offers a promising advancement, particularly suited for applications requiring long-context understanding.

Despite its promise, this TTT model is relatively untested across diverse domains due to its recent release, with only 28 citations as of now. We have an opportunity to both validate and potentially expand the model's utility in complex, real-world applications.

To establish baseline performance and enhance the model's robustness, we incorporated the Project Gutenberg (PG-19) dataset into our pre-training and evaluation pipeline. The PG-19 dataset, composed of over 28,000 books from Project Gutenberg. This makes for an ideal baseline for evaluating the TTT model's capacity to generalize across varying text domains.

Our primary objective is to evaluate the TTT model's adaptability through fine-tuning on domain-specific datasets, specifically the PubMed Medical Dataset and GovReport.

By fine-tuning TTT on these specialized datasets, we aim to not only validate its robustness and adaptability across distinct linguistic domains but also explore the broader applicability of long-context RNN architectures for specialized, high-stakes fields like medicine, public policy, and beyond.

## 2 Context

Sun et al. evaluate their model across only the Pile and Books (a subset of Pile) datasets. So, by training the TTT model architecture on domain-specific datasets, we are evaluating the model capability on capturing large scale datasets with more infrequent words. We continue Sun et al.'s choice of task of summarization because the TTT model is designed to take in (i.e. memorize/capture) large amounts of text.

As said before, the TTT model architecture leverages the strengths of both Transformers and RNNs, enabling it to process large-scale datasets efficiently and effectively. The summarization task involves condensing a longer piece of text into a shorter version while preserving the main ideas and information.

35 Other researchers have modified the TTT architecture further and evaluated its performance on other  
36 datasets. Xu et al. [2] have combined TTT layers with visual backbone layers (such as fourier  
37 transforms and convolutional layers). For the purposes of our project, we avoid making adjustments  
38 to the model architecture, but wanted to note that the TTT architecture is beginning to gain such  
39 traction.

40 The PubMed dataset, comprises biomedical research articles and their abstracts. Previous models  
41 trained on the PubMed dataset have performances shown in [3], with the top-performing model being  
42 the Top Down Transformer (AdaPool). AdaPool excels at leveraging hierarchical representations for  
43 summarization tasks by dynamically pooling information across layers to focus on salient content.

44 The GovReport dataset, consists of lengthy government reports and their summaries. Previous  
45 models trained on GovReport dataset have performances reported in [4]. With significantly fewer  
46 testing, BART has so far demonstrated the best performance. BART employs a Transformer-based  
47 encoder-decoder architecture pre-trained using a denoising autoencoder approach to reconstruct  
48 corrupted inputs. These results provide benchmarks for comparison as we extend the evaluation of  
49 the TTT model architecture to these domain-specific datasets.

### 50 3 Method

51 Our experiments involve loading the pre-trained TTT model from the Hugging Face Model Hub,  
52 leveraging the model’s existing training on the Pile dataset as a foundation. This initial step provides  
53 a solid baseline, allowing us to measure the impact of domain-specific fine-tuning on specialized  
54 datasets.

55 Using the Hugging Face Transformers library, we fine-tune the TTT model on our selected  
56 datasets—PubMed Medical Dataset and GovReport. We monitor the loss during training and  
57 report perplexity before and after fine-tuning, which are the same metrics reported in the original TTT  
58 paper. To ensure computational feasibility, we use the smallest pre-trained model available, the 125M  
59 parameter TTT-Linear, which has a linear layer as its hidden state. This architecture was chosen  
60 for its manageable size and balanced trade-off between representational capacity and efficiency.  
61 Fine-tuning experiments are conducted until performance converges, with periodic evaluations using  
62 the designated metrics to track progress.

63 Additionally, we explore an alternative fine-tuning approach using Low-Rank Adaptation (LoRA).  
64 LoRA enables parameter-efficient fine-tuning by introducing trainable low-rank matrices into the  
65 model’s architecture, significantly reducing the number of trainable parameters while preserving  
66 the model’s capacity for learning domain-specific nuances. This approach is particularly useful for  
67 fine-tuning larger models or scenarios with limited computational resources. By implementing LoRA,  
68 we aim to evaluate whether this method can achieve comparable or even superior results to standard  
69 fine-tuning, offering a scalable solution for handling larger model variants in future work.

70 To provide a broader context for our findings, we also establish perplexity baselines for the PubMed  
71 and GovReport datasets using Google’s T5 (Text-to-Text Transfer Transformer) model. T5, a state-  
72 of-the-art language model, was fine-tuned on the same datasets to enable direct comparison with the  
73 TTT model’s performance. By leveraging T5’s versatile architecture and pretraining on a diverse  
74 corpus, we aim to determine how TTT’s performance compares to that of a highly competitive SOTA  
75 model. These additional experiments provide valuable insights into the strengths and weaknesses of  
76 the TTT architecture and help benchmark its effectiveness in domain-specific tasks.

77 The fine-tuning objective remains consistent with the TTT model’s original pre-training setup—next-  
78 token prediction. This ensures that observed performance improvements can be directly attributed  
79 to fine-tuning rather than differences in training objectives. As for hyperparameters, we use the  
80 default learning rate of  $5e-5$ , a batch size of 4, and train for one epoch, meaning the model sees each  
81 training example once. Training is conducted on a single A100 GPU via Google Colab, providing  
82 sufficient computational resources to complete the experiments efficiently. This combination of  
83 standard fine-tuning, LoRA implementation, and SOTA model benchmarking allows us to rigorously  
84 evaluate the TTT model’s adaptability and utility in specialized domains.

## 85 4 Setup

86 This project involved a series of experiments designed to evaluate the Test-Time Training (TTT)  
 87 model’s ability to adapt to domain-specific datasets through fine-tuning and other optimization  
 88 techniques. All experiments were testing with a test set of 100 samples and training set of 17,917  
 89 samples for Gov Reports, and 19,200 samples for PubMed and Gutenberg. The experiments began  
 90 with a baseline evaluation, where the pre-trained TTT model, loaded from the Hugging Face Model  
 91 Hub, was tested without fine-tuning to establish a reference point for perplexity. This baseline  
 92 assessment provided a foundation for comparing the effects of subsequent fine-tuning and pre-training  
 93 methods.

94 The next experiment involved fine-tuning the TTT model on domain-specific datasets using a standard  
 95 configuration, including a learning rate of  $5e-5$ , a per device train batch size of 4, and one training  
 96 epoch or a max of 4800 steps, in alignment with the TTT paper. This step aimed to assess how  
 97 effectively the model could adapt to the linguistic and structural nuances of specialized content,  
 98 improving its ability to predict sequences accurately. During fine-tuning, the cross-entropy loss  
 99 function was employed to optimize the model’s next-token predictions. Additionally, Low-Rank  
 100 Adaptation (LoRA), a parameter-efficient fine-tuning technique, was tested to reduce computational  
 101 overhead by modifying only a subset of the model’s parameters.

102 To examine the impact of pre-training, another experiment incorporated the Gutenberg dataset as  
 103 an intermediate step before fine-tuning. This experiment was designed to test whether additional  
 104 pre-training on a large corpus of text with long sequences could enhance the TTT model’s ability to  
 105 handle specialized tasks during fine-tuning. While the original TTT paper used the Books dataset for  
 106 pre-training, it is no longer available due to copyright restrictions, making the Gutenberg dataset a  
 107 suitable alternative for this experiment.

108 In another experiment, pre-training on the Gutenberg dataset was combined with LoRA-based fine-  
 109 tuning. This combination aimed to leverage the benefits of both approaches—enhanced contextual  
 110 understanding from pre-training and the efficiency of LoRA fine-tuning. Together, these experiments  
 111 provided a comprehensive framework for evaluating the adaptability, efficiency, and potential of the  
 112 TTT model across various fine-tuning strategies.

113 Finally, a benchmarking experiment was conducted using T5-220M, a transformer-based language  
 114 model. The T5-220M model was used to establish baseline perplexities for comparison against the  
 115 TTT model’s performance, providing a reference point for understanding the strengths and limitations  
 116 of the smaller TTT model in relation to a larger, state-of-the-art architecture.

## 117 5 Outcomes and Results

	Gutenberg	PubMed	GovReport
TTT-125M	36473	37366	37422
Fine-Tuned	40	171	225
Fine-Tuned w/ LoRA	-	14356	15361
Fine-Tuned (Gutenberg)	40	3948	4513
Fine-Tuned (Gutenberg) & Fine-tuned w/ LoRA	-	750	660
T5-220M	-	2194176	56505511

Table 1: Perplexity across models and datasets. The base model is TTT-125M. Fine-tuned means fine-tuned on the dataset it is evaluated on.

### 118 5.1 Baseline (T5)

119 It is remarkable how much better the base TTT-125M model performs compared to the T5-220M  
 120 model. The T5 is one of the state-of-the-art (SOTA) language models, developed at Google, and at  
 121 220M parameters, is almost twice the size of the TTT model version we test against. We are sure  
 122 that if we were to fine-tune the T5 model on the respective datasets that the performance would  
 123 significantly improve, as it has for the TTT-125M model. It is still useful to compare the base  
 124 (non-fine-tuned) TTT model against the base T5 model; the original TTT authors do not compare the  
 125 TTT model against popular SOTA models.

## 126 5.2 Loss




Figure 1: PubMed Loss Graph




Figure 2: GovReport Loss Graph

127 Above are the loss graphs for the TTT-125M model fine-tuned and evaluated on PubMed and  
 128 GovReport respectively. The fine-tuning on the PubMed data shows a gradual decrease in loss over  
 129 the course of the epoch, with the model starting with a relatively high loss of 6.03 at step 500 and  
 130 decreasing to 4.6 at step 2500; the model learns to fit the train data better. The fine-tuning on the  
 131 GovReport dataset shows steady improvement throughout the training process, with the loss gradually  
 132 decreasing from 6.58 at step 500 to 4.43 at step 4000, reflecting the model’s increasing ability to  
 133 adapt to the specific characteristics of the new dataset.

## 134 5.3 Perplexity

135 We choose to use perplexity as the metric for performance to follow the methodology of Sun et al.  
 136 Perplexity is a metric commonly used to evaluate the quality of language models. It measures how  
 137 well a probability distribution or probability model predicts a sample. It is defined as the inverse  
 138 probability of the test set, normalized by the number of words (or tokens) in the text. Formally,

$$\text{Perplexity} = \exp \left( -\frac{1}{N} \sum_{i=1}^N \log P(w_i) \right)$$

139 where:

- 140 •  $N$  is the total number of words in the sequence.
- 141 •  $w_i$  represents the  $i$ -th word in the sequence.
- 142 •  $P(w_i)$  is the probability assigned by the model to the word  $w_i$ .

143 Lower perplexity values indicate that the model assigns higher probabilities to the actual sequence of  
 144 words, and thus, it performs better.

### 145 5.3.1 Fine-Tuning on Respective Datasets

146 The perplexity values before and after fine-tuning reveal a significant improvement in the model’s  
 147 performance: it drops from an extremely high 37366.53 to a much more reasonable 171.17 after  
 148 fine-tuning on the medical dataset, indicating that the fine-tuned model is far superior at summarizing  
 149 medical-related text.

150 The perplexity results reveal a notable improvement in model performance: before fine-tuning, the  
 151 perplexity was exceedingly high at 37422.87, but after training, it dropped substantially to 225.66,  
 152 indicating that the fine-tuned model is far superior at summarizing the government reports.

153 It does not surprise us that fine-tuning on the respective domain-specific training datasets before  
 154 testing on their evaluation datasets provides the strongest performance. This aligns with established  
 155 principles in transfer learning, where adapting a pre-trained model to the distribution and linguistic  
 156 characteristics of a target domain significantly improves downstream task performance. The observed  
 157 reduction in perplexity underscores the model’s enhanced ability to generate domain-relevant and  
 158 semantically coherent summaries. These findings emphasize the effectiveness of domain-specific

fine-tuning in narrowing the representation gap between pre-trained models and specialized datasets, particularly for tasks requiring a deep understanding of technical or specialized language. See Appendix for best govReport output examples from fine-tuning.

### 5.3.2 LoRA

LoRA (Low-Rank Adaptation) is a parameter-efficient fine-tuning technique for large language models (LLMs) that reduces memory and computational costs. Instead of updating all the model's weights, LoRA keeps the pre-trained weights frozen and injects small trainable low-rank matrices into specific layers, such as attention mechanisms.

We apply LoRA using the `PeftConfig()` library from Hugging Face. Our hyperparameters for this were  $r = 8$ ,  $\alpha = 32$ , and dropout rate = .1, where  $r$  represents the rank of the low-rank decomposition matrices used in LoRA,  $\alpha$  is the scaling factor that adjusts the contribution of the low-rank updates, and the dropout rate controls the fraction of neurons randomly set to zero during training to prevent overfitting.

We experimented with LoRA to assess the tradeoff between performance and efficiency. The results, shown in Table 1, reveal that while fine-tuning the TTT-125M model without LoRA resulted in low perplexity values (40 for Gutenberg, 171 for PubMed, and 225 for GovReport), applying LoRA increased perplexity, especially for PubMed and GovReport (14356 and 15361, respectively). This indicates that LoRA, while reducing computational overhead, leads to higher perplexity and may not perform as well as full fine-tuning in certain datasets. Additionally, fine-tuning with LoRA reduced the number of trainable parameters from 110,276,688 to 589,824 and lowered GPU memory usage from 36% to 26%. Using these hyper parameters, only 0.53% of the 125M parameters are trained. It is unsurprising that the performance is significantly worse, because the capacity of the model has clearly significantly decreased. Future experiments may seek to evaluate the performance of fine-tuning using LoRA on larger model version of TTT (more than 1 Billion), given that the main benefit of LoRA is to improve training efficiency, and it is more important to train efficiently for larger models.

### 5.3.3 Fine-Tuning only on Gutenberg

Fine-tuning on the Gutenberg dataset was chosen as a substitute for the Books dataset, which was used in the original TTT paper to train the TTT-125M model. However, due to copyright issues, the Books dataset is no longer available for public use. Given this limitation, the Gutenberg dataset, which offers a large collection of publicly available literary texts, was selected as a suitable alternative.

When LoRA was applied to the model fine-tuned on the Gutenberg dataset, the perplexity for PubMed and GovReport remained competitive (750 and 660, respectively) compared to models fine-tuned directly on those datasets (PubMed: 171, GovReport: 225). This suggests that the model, having learned generalizable language patterns from Gutenberg, could transfer some of its capabilities to other domains. In contrast, models fine-tuned specifically on PubMed and GovReport performed worse when LoRA was applied, with perplexity values increasing significantly (PubMed: 14356, GovReport: 15361). This highlights that while LoRA reduces computational overhead and memory usage, it may also limit the model's ability to adapt to specialized content compared to full fine-tuning.

In summary, fine-tuning on the Gutenberg dataset allowed the TTT-125M model to perform well across multiple datasets, especially with LoRA applied. The model's ability to transfer knowledge from Gutenberg to domains like PubMed and GovReport highlights the value of a diverse fine-tuning corpus. While LoRA reduced memory usage and computational overhead, it introduced a performance tradeoff, particularly for more specialized datasets.

## 6 Future Work

This project demonstrates the effectiveness of fine-tuning the TTT model on domain-specific datasets, significantly improving perplexity for summarization tasks in specialized fields like biomedical literature and government reports. We further experiment with fine-tuning using PEFT techniques like LoRA and only fine-tuning on a more general purpose dataset like Project Gutenberg; as expected, both of these techniques reduce performance (measured by perplexity) compared with traditional fine-tuning. We proved the utility of the TTT model and establish its effectiveness across a broader range of datasets.

## 7 Github Repository Link

<https://github.com/anyaeross18/ttt-AML>.

## References

- [1] Hugging Face. *Training and Evaluation with the Transformers Library*. Hugging Face, 2024. Web. Accessed 9 Nov. 2024. <https://huggingface.co/docs/transformers/training#evaluate>.
- [2] Sun, Yu, et al. *Learning to (Learn at Test Time): RNNs with Expressive Hidden States*. *arXiv*, 2024, arXiv:2404.06252. Web. Accessed 9 Nov. 2024. <https://arxiv.org/abs/2404.06252>.
- [3] “Text Summarization on PubMed.” *Papers with Code*, Papers with Code. Accessed 13 Dec. 2024. <https://paperswithcode.com/sota/text-summarization-on-pubmed-1>.
- [4] “Text Summarization on GovReport.” *Papers with Code*, Papers with Code. Accessed 13 Dec. 2024. <https://paperswithcode.com/sota/text-summarization-on-govreport>.

## 8 Appendix

This section presents the generated output from the TTT-125m model, fine-tuned on the Gov Reports dataset. The model was evaluated based on its ability to produce a summary for a given report. The outputs displayed here demonstrate the results for the first report in the test split of the Gov Reports dataset, showing both the base model output (which includes random characters) and the output from the model that achieved the best perplexity.

### 8.1 Input: First Report in the Test Split of the Gov Reports Dataset

*Note: Only the first 3000 characters are shown.*

In our prior work, we have found that technological innovation involves not only creating new ideas but also translating those ideas into a new product or service. Innovation, and the research driving it, is inherently risky because the likelihood that research can be translated into a product or service and the ultimate value of that product or service are unknown. The Department of Commerce’s National Institute of Standards and Technology describes the path from innovation to commercialization as comprised of three overarching stages: inventing, transitioning to making, and selling. (See fig. 1 for a description of the path from innovation to commercialization.) FDA and USDA have responsibility for overseeing the safety of the food supply. In general, FDA is responsible for ensuring the safety of virtually all domestic and imported food products except those regulated by USDA. USDA is responsible for ensuring the safety of meat, poultry, processed egg products, and catfish. FDA and USDA cooperate with states, tribes, and local food safety and public health agencies to carry out their federal responsibilities. FDA and USDA carry out their responsibilities in part through inspections of facilities where food is produced. The frequency of inspections the agencies conduct varies, as follows: FDA. FDA’s authority requires a risk-based approach, in which inspection rates vary depending on the level of risk associated with a food product. FDA conducts risk-based inspections of high-risk and non-high-risk food facilities. For example, the FDA Food Safety Modernization Act, signed into law in 2011, specified that FDA had to inspect all high-risk domestic facilities at least every 3 years. USDA. Depending on the type of facility, USDA conducts inspections at least once per operating shift or maintains a constant presence. Specifically, USDA conducts carcass-by-carcass inspection at all federally inspected meat and poultry slaughter facilities and verifies that these establishments follow all food safety and humane handling requirements. At facilities that process meat and poultry products, USDA conducts inspections at least once per production shift, following the agency’s longstanding interpretation of its statutes requiring it to do so. Among other things, the Federal Food, Drug, and Cosmetic Act requires that food additives be approved by FDA before they can be lawfully used in foods. Substances added

260 to food are considered unsafe unless the agency establishes that the use of the food  
261 additive, under specific conditions for use, will be safe, or unless the substance  
262 is generally recognized as safe (GRAS) under the conditions of its intended use  
263 among qualified experts. As we reported in 2010, the Federal Food, Drug, and  
264 Cosmetic Act exempts GRAS substances from the act's general requirement that  
265 companies obtain FDA approval before marketing food containing a new additive.  
266 GRAS substances include hundreds of spices and artificial flavors, emulsifiers

## 267 8.2 Summary: First Report Summary in the Gov Reports Dataset

268 Multiple firms have produced cell-cultured meat as part of their research and devel-  
269 opment. These products appear likely to become available to consumers in coming  
270 years. FDA and USDA are the primary agencies responsible for overseeing the  
271 safety of the nation's food supply. However, some stakeholders have expressed  
272 concern about the agencies' oversight of cell-cultured meat amidst a fragmented  
273 federal food safety oversight system. GAO was asked to review federal oversight  
274 of cell-cultured meat. This report (1) describes what is known about methods for  
275 commercially producing cell-cultured meat, and (2) examines the extent to which  
276 FDA and USDA are collaborating to provide regulatory oversight of cell-cultured  
277 meat. GAO conducted a literature review; reviewed documentation from FDA,  
278 USDA, and stakeholder groups; analyzed public comments submitted to the agen-  
279 cies; compared agency efforts with leading practices for interagency collaboration;  
280 and conducted site visits to selected cell-cultured meat firms. General information  
281 about the process of making cell-cultured meat—food products grown from the  
282 cells of livestock, poultry, and seafood—is available. However, no company is com-  
283 mercially producing cell-cultured meat. Specific information about the technology  
284 being used, eventual commercial production methods, and composition of the final  
285 products is not yet known. The general process contains five phases: biopsy, cell  
286 banking, growth, harvest, and food processing (see figure). The technology and  
287 methods to be used for commercial production are still in development, and produc-  
288 ers, regulators, and consumers do not have clarity about many specifics about the  
289 process and final product. For example, it is unclear whether production methods  
290 and products will use or contain genetically-engineered cells or medications such  
291 as antibiotics. The Food and Drug Administration (FDA) and U.S. Department  
292 of Agriculture (USDA) have begun collaborating on regulatory oversight of cell-  
293 cultured meat. For example, in 2019, the agencies signed an interagency agreement  
294 and created three working groups to carry out the terms of the agreement. How-  
295 ever, the agreement and working groups could more fully incorporate practices  
296 to enhance and sustain collaboration, such as defining outcomes. For example,  
297 the agreement identifies the development of labeling principles as an outcome,  
298 but does not describe how the agencies will track and monitor progress toward  
299 this outcome, and the working groups identify a lead agency but not members'  
300 roles. Also, agency officials said they decided FDA would oversee cell-cultured  
301 seafood other than catfish, but they have not formally announced or documented  
302 this decision. Developing and updating written guidance and agreements is also a  
303 leading practice for interagency collaboration. By fully incorporating leading prac-  
304 tices into their efforts to collaborate, the agencies could minimize potential overlap  
305 and fragmentation, use resources in a more efficient manner, and better ensure the  
306 public and other key stakeholders have clarity about the agencies' oversight respon-  
307 sibilities. GAO recommends that FDA and USDA more fully incorporate leading  
308 practices for effective collaboration in the agencies' interagency agreement. FDA  
309 and USDA partially concurred and indicated a willingness to incorporate these  
310 practices in a more detailed agreement, which would also meet the intent of the  
311 recommendations. The agencies concurred with the four other recommendations.

## 312 8.3 Base Model Output (Random Characters)

313 The base model output consists of randomly generated characters, highlighting the model's inability  
314 to generate coherent text prior to fine-tuning. Below is the output produced by the base model.

315           rix()‘,uteur JoseLayout conditionpont valuesospelSESSION dentopacity lets lov  
 316           macroemhab waste→ <? Atlas traversdomain Pok proces restouchDevice libraries-  
 317           rame basketballon seulement trois seculo varietyElementiras hal mucount

#### 318 **8.4 Best Model Output (Lowest Perplexity)**

319   The output below was generated by the fine-tuned model that achieved the best perplexity on the Gov  
 320   Reports dataset. This summary reflects the model’s ability to process the input report and generate a  
 321   more coherent and relevant summary.

322   *Note: The first half of the model’s output matches the input report exactly.*

323           In our prior work, we have found that technological innovation involves not only  
 324           creating new ideas but also translating those ideas into a new product or service.  
 325           Innovation, and the research driving it, is inherently risky because the likelihood  
 326           that research can be translated into a product or service and the ultimate value of  
 327           that product or service are unknown. The Department of Commerce’s National  
 328           Institute of Standards and Technology describes the path from innovation to com-  
 329           mercialization as comprised of three overarching stages: inventing, transitioning  
 330           to making, and selling. (See fig. 1 for a description of the path from innovation  
 331           to commercialization.) FDA and USDA have responsibility for overseeing the  
 332           safety of the food supply. In general, FDA is responsible for ensuring the safety  
 333           of virtually all domestic and imported food products except those regulated by  
 334           USDA. USDA is responsible for ensuring the safety of meat, poultry, processed  
 335           egg products, and catfish. FDA and USDA cooperate with states, tribes, and local  
 336           food safety and public health agencies to carry out their federal responsibilities.  
 337           FDA and USDA carry out their responsibilities in part through inspections of facil-  
 338           ities where food is produced. The frequency of inspections the agencies conduct  
 339           varies, as follows: FDA. FDA’s authority requires a risk-based approach, in which  
 340           inspection rates vary depending on the level of risk associated with a food product.  
 341           FDA conducts risk-based inspections of high-risk and non-high-risk food facilities.  
 342           For example, the FDA Food Safety Modernization Act, signed into law in 2011,  
 343           specified that FDA had to inspect all high-risk domestic facilities at least every 3  
 344           years. USDA. Depending on the type of facility, USDA conducts inspections at  
 345           least once per operating shift or maintains a constant presence. Specifically, USDA  
 346           conducts carcass-by-carcass inspection at all federally inspected meat and poultry  
 347           slaughter facilities and verifies that these establishments follow all food safety  
 348           and humane handling requirements. At facilities that process meat and poultry  
 349           products, USDA conducts inspections at least once per production shift, following  
 350           the agency’s longstanding interpretation of its statutes requiring it to do so. Among  
 351           other things, the Federal Reserve has been used to provide information about an  
 352           interagency officials’s financial institutions that include private sector services.  
 353           To determine whether the number of this report will help develop both access  
 354           activities such as they may use some of environmental operations; assessment  
 355           policies and procedures—such as well as “aid” projects—income individuals who  
 356           are unemployment in figure 1). While FEMA must meet the development of time  
 357           of their own needs of the contractor’ ability to obtain, including the CMS and how  
 358           to achieve government employees, among others as the program, the agency, and  
 359           funded. These issues require agency management systems to identify and manage  
 360           the same cost-term care system. This plan managing the primary costs and